

Design Consistent Hashing

- To achieve horizontal scaling, it is important to distribute requests / data efficiently and evenly across servers. Consistent hashing is a commonly used technique to achieve this goal.
- The Rehashing Problem
 - If you have n cache servers, a common way to balance the load is to use following hash method:
 $\text{serverIndex} = \text{hash}(\text{key}) \% N$, N -Size of server pool

key	hash	hash % 4
key0	18358617	1
key1	26143584	0
key2	18131146	2
key3	35863496	0
key4	34085809	1
key5	27581703	3
key6	38164978	2
key7	22530351	3

- This approach works well when the size of server pool is fixed, and data distribution is even. However, problem arises when servers are added or removed.
- Consistent Hashing
 - Consistent hashing is a special kind of hashing such that when a hash table is re-sized and consistent hashing is used, only k/n keys need to be remapped on average, where k is number of keys and n is the number of slots. In

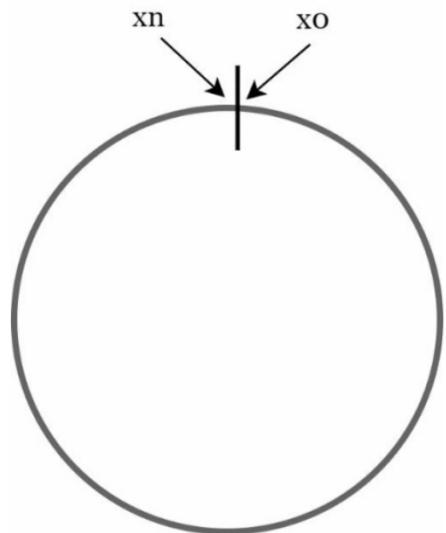
contrast, most traditional hash tables, a change in number of array slots causes nearly all keys to be remapped.

- Hash Space and Hash Ring

- Assume SHA-1 is used as the hash function f , and the output range of the hash function is: $x_0, x_1 \dots x_n$. SHA-1 hash space goes from 0 to $2^{160}-1$.

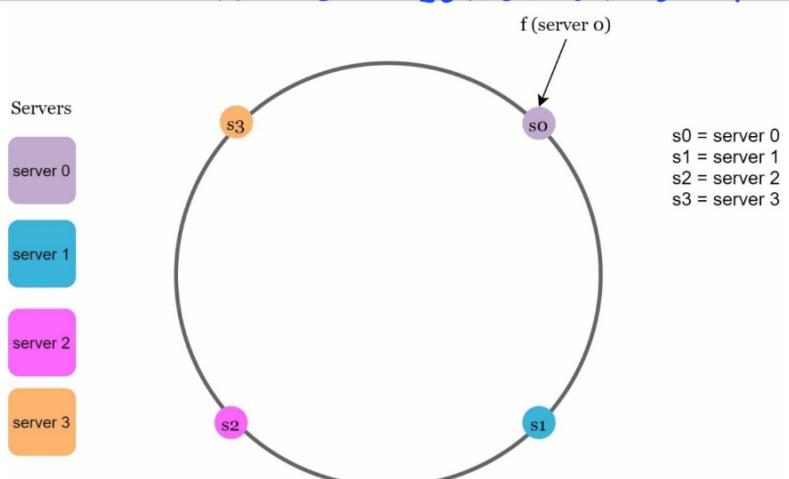


- By connecting both ends, we get a hash ring.



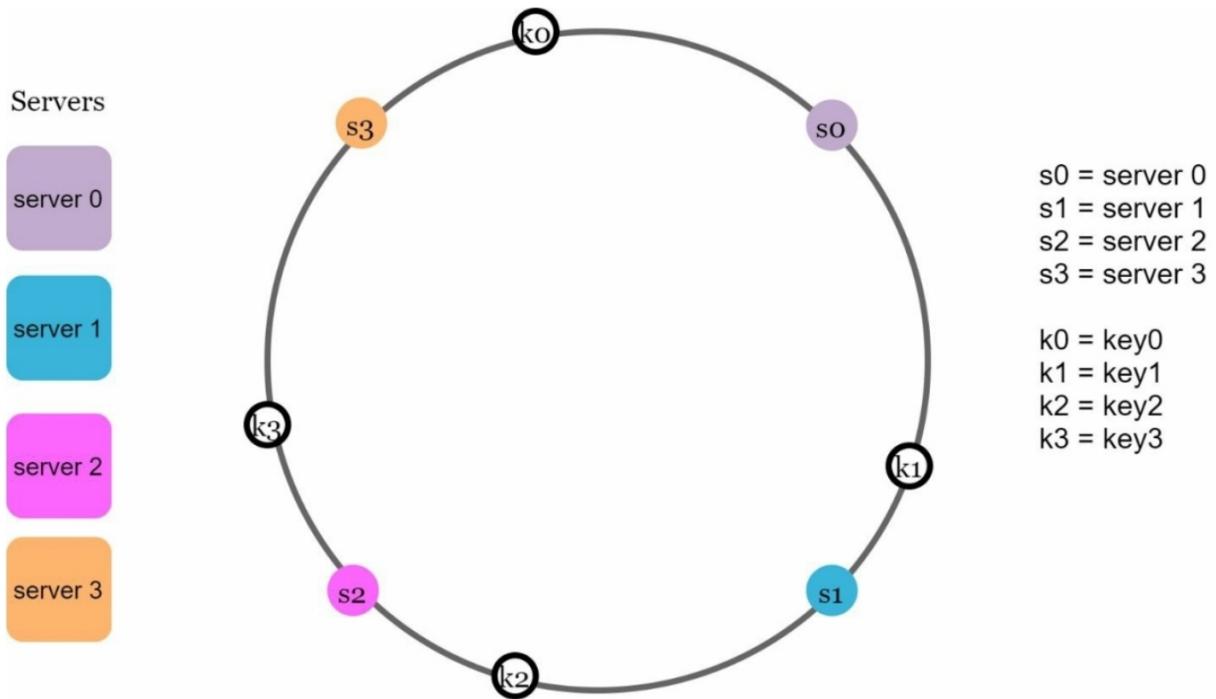
- Hash Servers

- Using the same hash function f , servers are mapped based on server IP or name onto the ring.



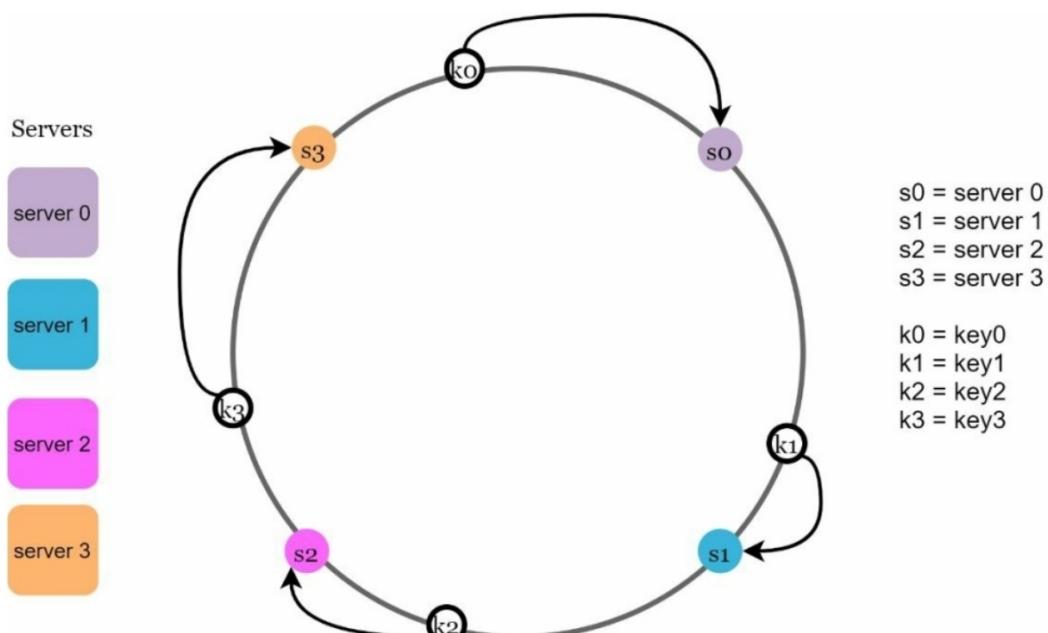
- Hash Keys

- Hash functions used here is different from the one in "the grehashing problem" and there is no modular operation.
- Cache keys are hashed onto the hash ring.



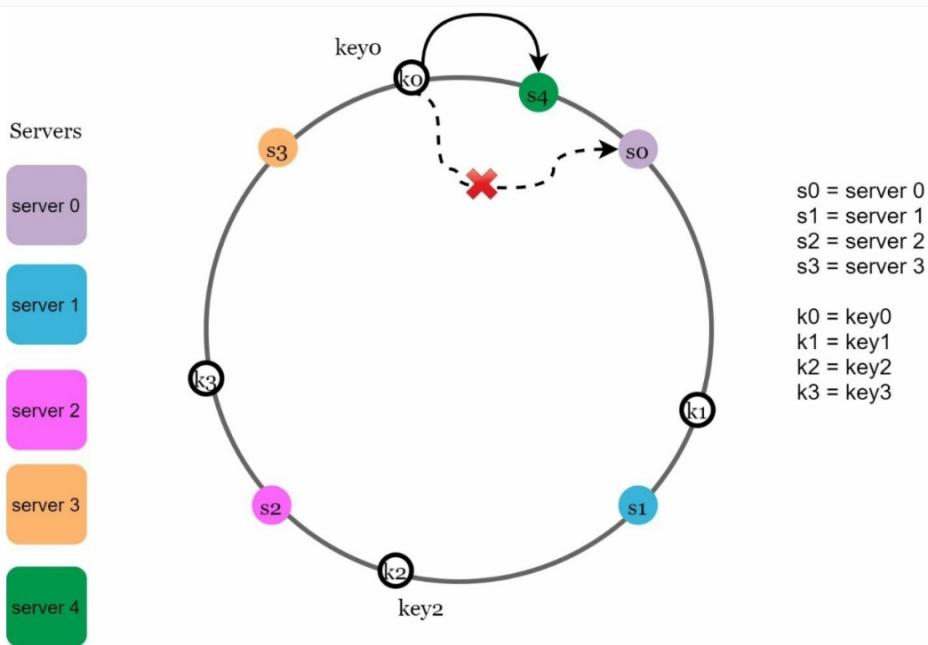
- Server Lookups

- To determine which server key is stored on, we go clockwise from the key position on the ring until a server is found.



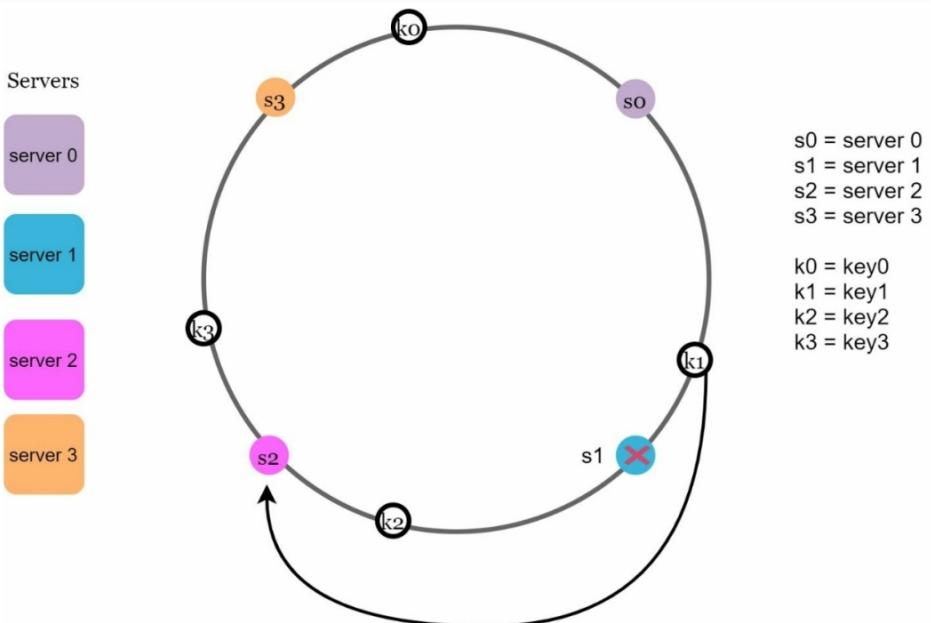
- Add a Server

- Adding a new server will only require redistribution of a fraction of keys.
- After a new server 4 is added, only key 0 needs to be redistributed. k1, k2, k3 remain on the same servers. Before server 4 is added, key 0 is stored on server 0. Now, key 0 will be stored on server 4 because server 4 is the first server it encounters by going clockwise from key 0's position on the ring.
- The other keys are not redistributed based on consistent hashing algorithms.

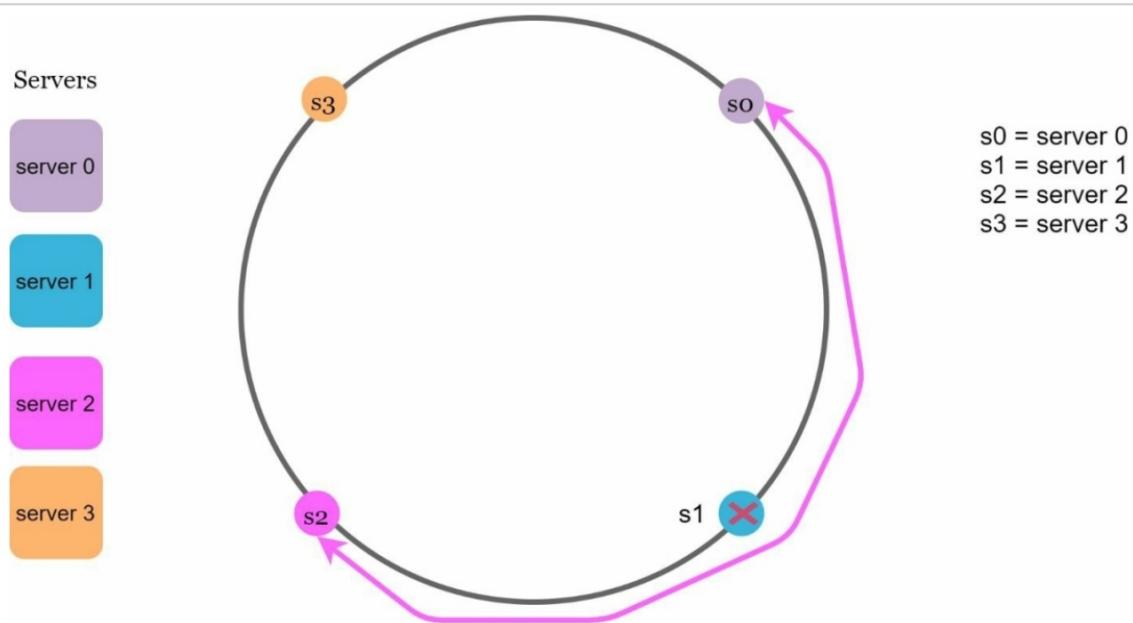


- Remove a Server

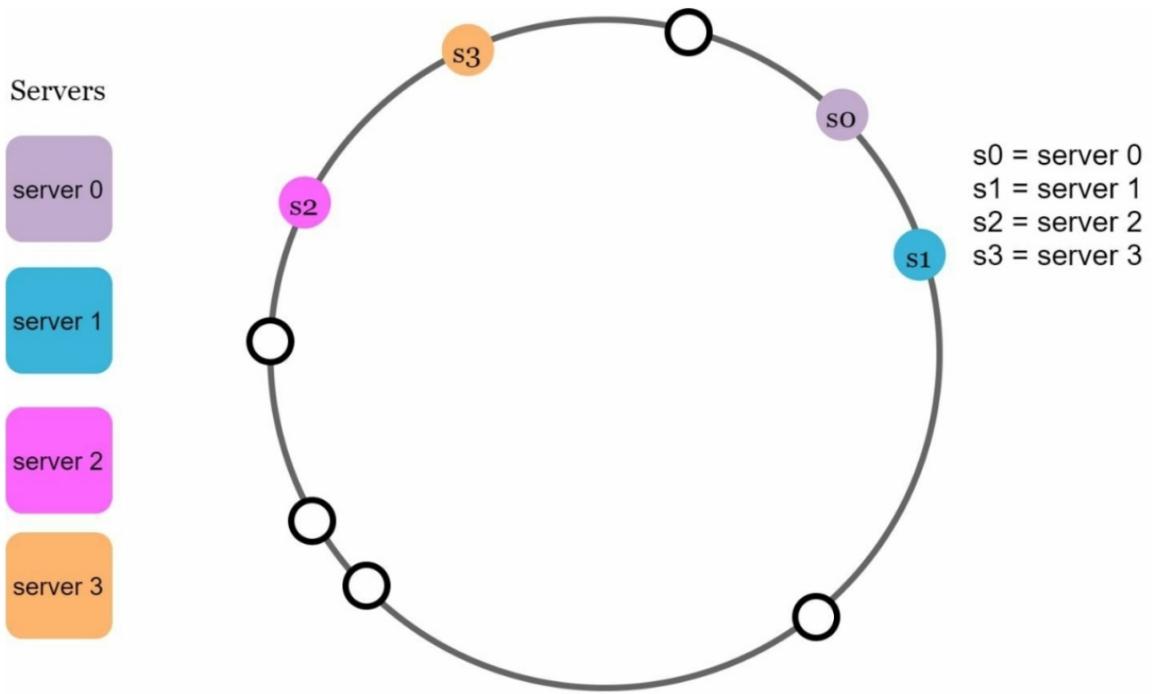
- When a server is removed, only a fraction of keys require redistribution with consistent hashing.
- When server 1 is removed, only key1 must be remapped to server 2.
- The rest of the keys are unaffected.



- Two issues in the basic approach
 - Basic steps for consistent hashing algorithms are :
 - Map servers and keys onto the ring using a uniformly distributed hash function.
 - To find out which server a key is mapped to, go clockwise from the key position until the first server on the ring is found.
 - First, it is impossible to keep the same size of partitions on the ring for all servers considering a server can be added or removed.

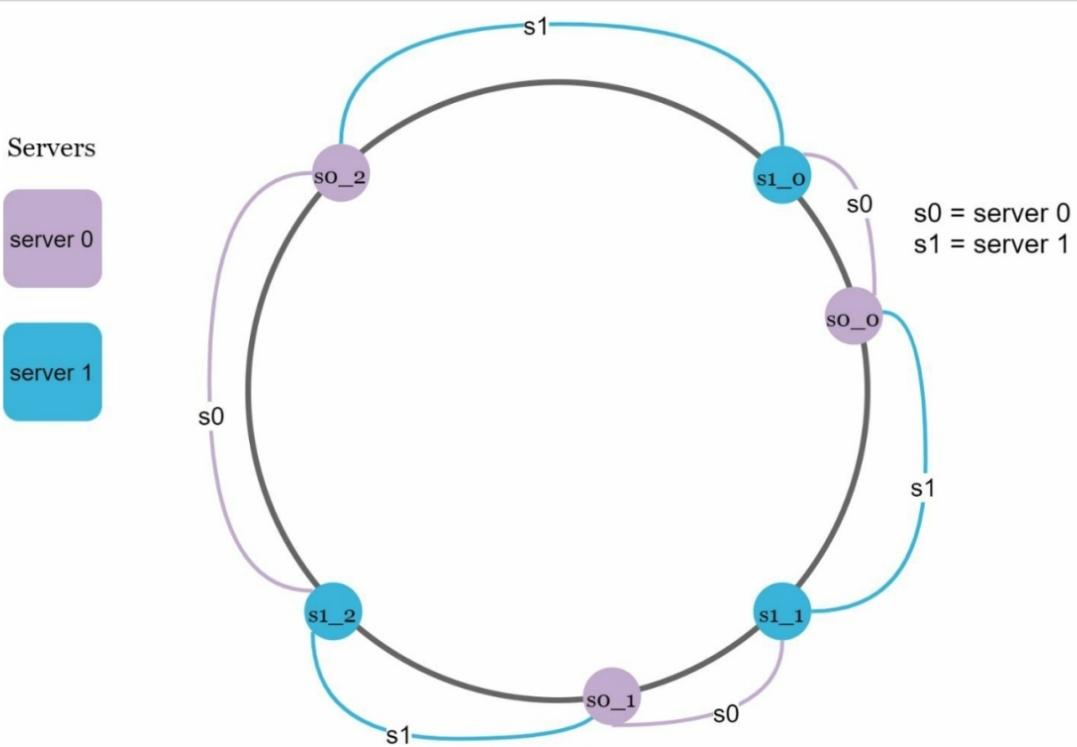


- Secondly, it is possible to have a non-uniform key distribution on the ring.

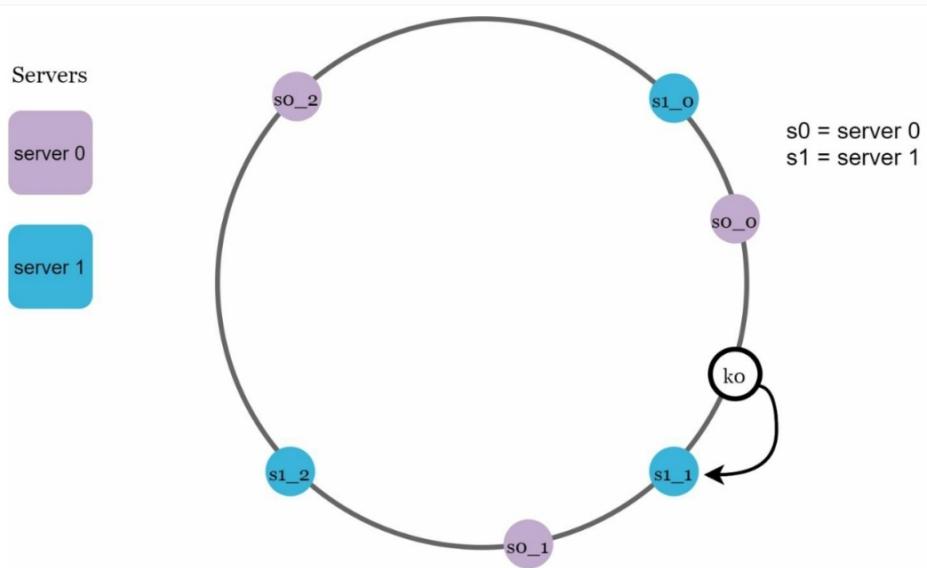


- Virtual Nodes

- A virtual node refers to the real node, and each server is represented by multiple virtual nodes on the ring.
- With virtual nodes, each server is responsible for multiple partitions.



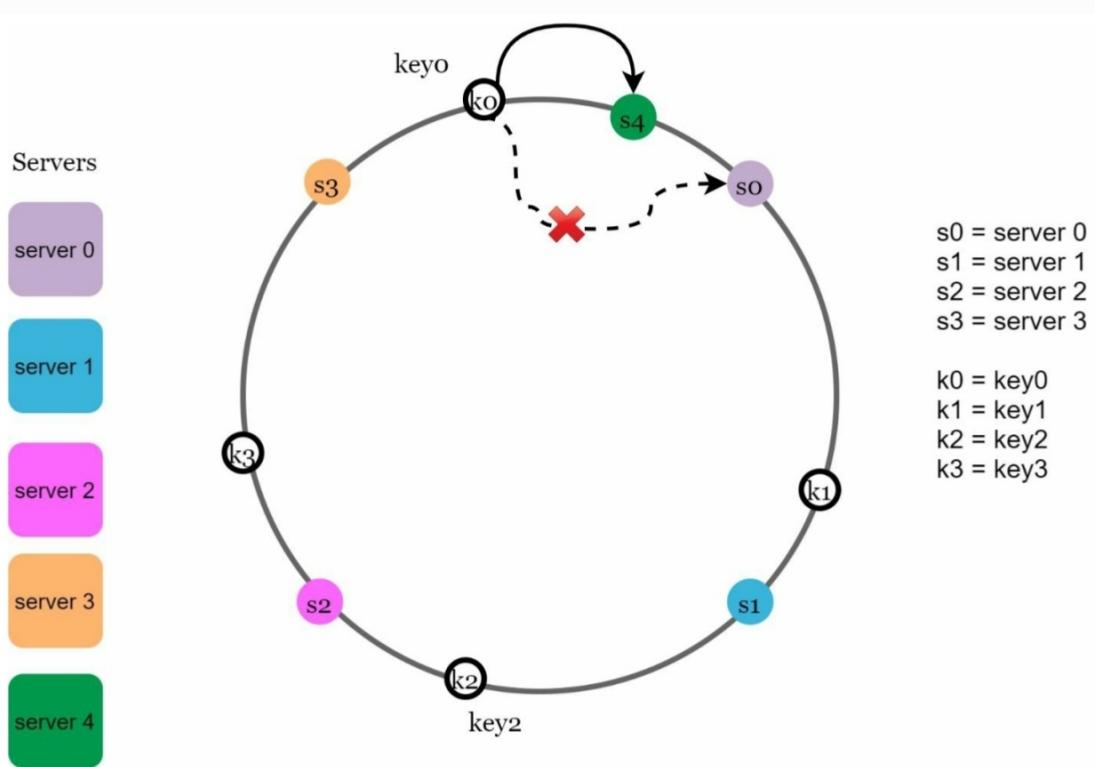
- To find which servers a key is stored on, we go clockwise from the key's location and find the first virtual node encountered on the ring.



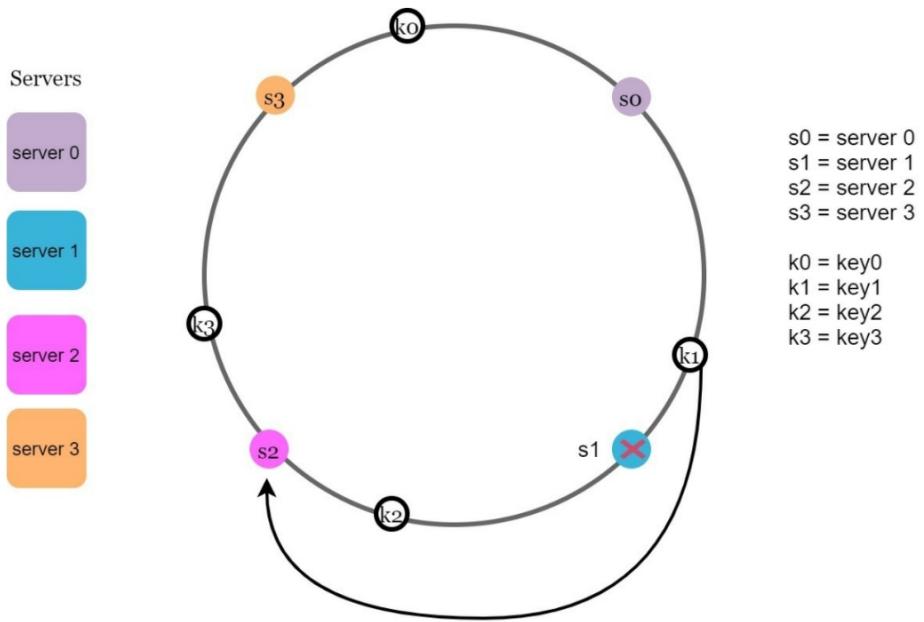
- As the number of virtual nodes increases, the distribution of keys becomes more balanced.

- Find Affected Keys

- When a server is added or removed, a fraction of data needs to be redistributed.



- Server 4 is added onto the ring. The affected range starts from s_4 and moves anticlockwise around the ring until a server is found. Thus, keys located between s_3 and s_4 need to be redistributed to s_4 .



- When a server is removed, the affected range starts from s_1 and moves anticlockwise around the ring until a server is found. Thus keys located between s_0 and s_1 needs to be redistributed to s_2 .
- Wrap Up
 - Benefits of consistent hashing
 - Minimized keys are redistributed when servers are added or removed.
 - It is easy to scale horizontally because data are more evenly distributed.
 - Mitigate hotspot key problem. Consistent hashing helps to mitigate the problem by distributing the data more evenly.