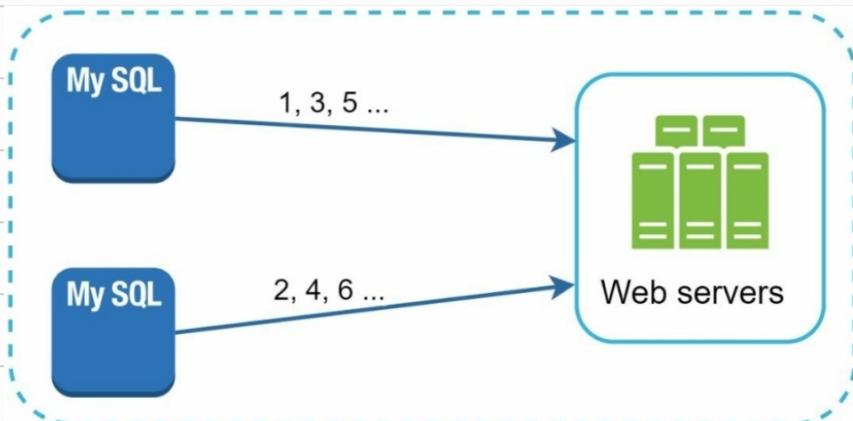


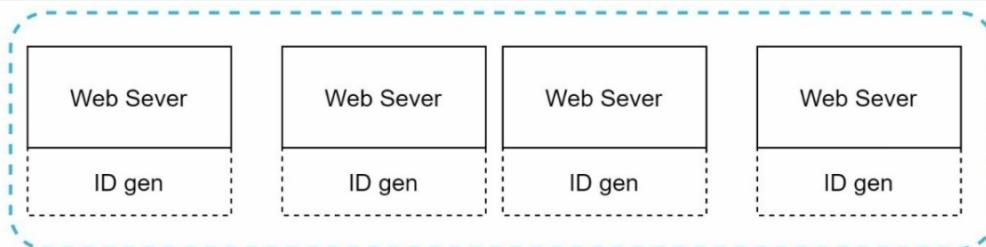
Design a Unique ID Generator in Distributed Systems

- Step 1 - Understand the problem and establish design scope
 - ID must be unique.
 - IDs are numerical value only.
 - IDs fit into 64-bit.
 - IDs are ordered by date.
 - Ability to generate over 10,000 unique per second.
- Step 2 - Propose high-level design and get buy-in
 - Multiple options can be used to generate unique IDs in distributed systems. The options are:
 - Multi-master replication
 - Universally Unique Identifiers (UUID)
 - Ticket Server
 - Twitter Snowflake Approach.
 - Multi-Master Replication



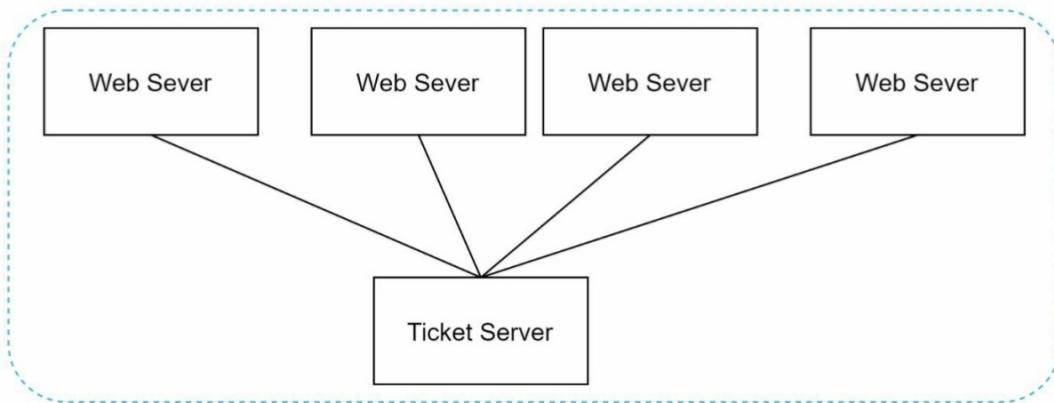
- This approach uses the database's auto-increment feature. Instead of increasing the next ID by 1, we increase it by k, where k is the number of database servers in use.

- This strategy has major drawbacks:
 - Hard to scale with multiple data centers.
 - IDs do not go up with time across multiple servers.
 - It does not scale when a server is added or removed.
- UUID
 - A UUUID is another way to obtain unique IDs. UUUID is a 128 bit number used to identify information in computer systems.
 - UUUID has a very low probability of getting collision.



- In this design, each web server contains an ID Generator, and a web server is responsible for generating IDs independently.
- Pros :
 - Generating UUUID is simple. No coordination between servers is needed so there will not be any synchronization issues.
 - The system is easy to scale because each web server is responsible for generating IDs they consume. ID generation can easily scale with web servers.
- Cons :
 - IDs are 128 bits long, but our requirement is 64 bits.
 - IDs do not go up with time.
 - IDs could be non-numeric.

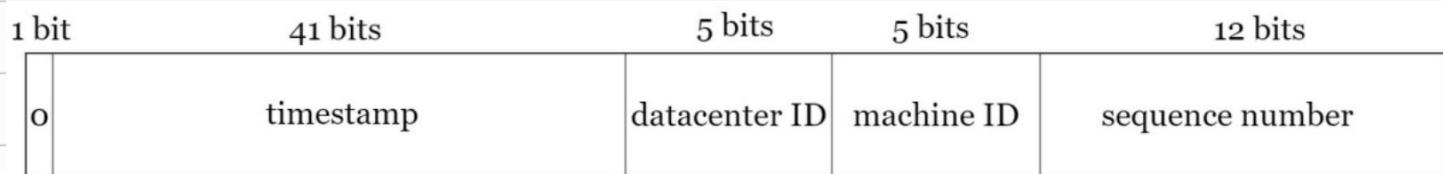
- Ticket Server
 - Tickets servers are another interesting way to generate unique IDs.
 - Flickr developed ticket servers to generate distributed primary keys.



- The idea is to use a centralized auto-increment feature in a single database server.
- Pros :
 - Numeric IDs.
 - It is easy to implement, and it works for small to medium-scale applications.
- Cons :
 - Single Point of Failure - Single ticket server means if the ticket server goes down, all systems that depend on it will face issues. To avoid a SPOF, we can set up multiple ticket servers. However, this will introduce new challenges such as data synchronization.

- Twitter Snowflake Approach

- Instead of generating an ID directly, we divide an ID into different sections.



• Each section is explained below:

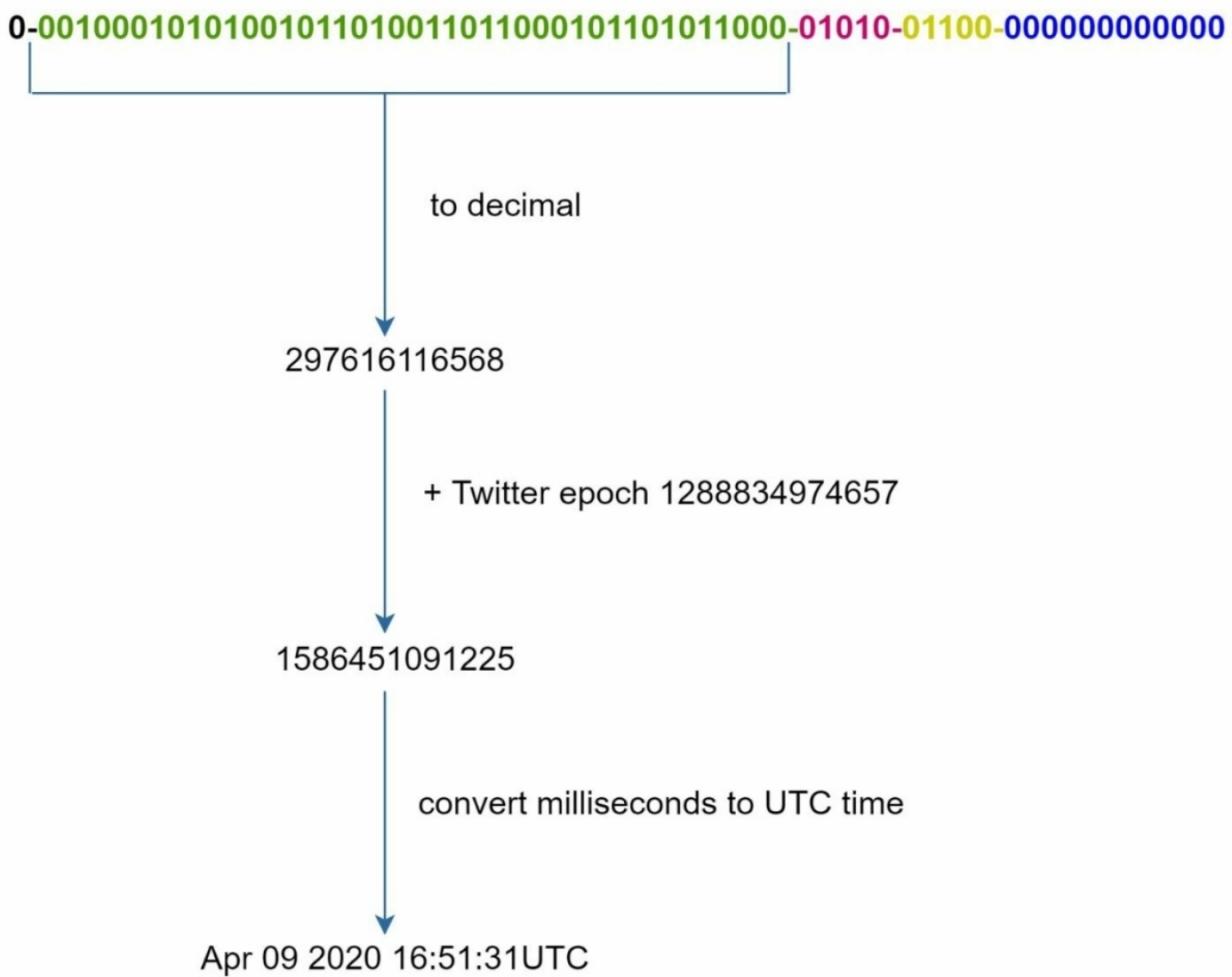
- Sign bit: 1 bit. It will always be 0. This is reserved for future uses. It can potentially be used to distinguish between signed and unsigned numbers.
- Timestamp: 41 bits. Milliseconds since the epoch or custom epoch.
- Datacenter ID: 5 bits, which gives us $2^5 = 32$ datacenters.
- Machine ID: 5 bits, which gives $2^5 = 32$ machines per datacenter.
- Sequence number: 12 bits, For every ID generated on that machine, the sequence number is incremented by 1. The number is reset to 0 every millisecond.

• Step 3 - Design Deep Dive

- Twitter Snowflake ID generator approach works best for the given requirements.
- Datacenter IDs and machine IDs are chosen at the startup time, generally fixed when the system is running.

- Timestamp

- The most important bits make up the timestamp section. As timestamps grow with time, IDs are sortable with time.



- The maximum timestamp that can be represented in 41 bits is $2^{41} - 1 = 219902325551$ ms ≈ 69 years. This means the ID generation will work for 69 years.
- Sequence Number
 - Sequence number is 12 bits, which gives 4096 combinations. This field is 0 unless more than one ID is generated in a ms. In theory, a machine can support a maximum of 4096 new IDs per ms.

- Step 4 - Wrap Up
 - Few additional talking points:
 - Clock Synchronization: In our design, we assume ID generation servers have the same clock. This assumption might not be true when a server is running on multiple cores or in multi-machine scenario. Network Time Protocol is the most popular solution to this problem.
 - Section length tuning: Fewer sequence numbers but more timestamp bits are effective for low concurrency and longterm applications.
 - High Availability: Since ID generator is a mission critical system, it must be highly available.